

Incoming image x_t



**Frozen detector-side
visual encoder**

(e.g., OpenCLIP ViT-B/16 or
Qwen3-VL-Embedding-2B)

Embedding $z_t = f(x_t) \in \mathbb{R}^d$



**Trained TAM-VLM
detection head**

Detector head and threshold
are *fixed after training*.

Detector score $s(x_t) \in [0, 1]$



Framewise decision

$$a_t = 1[s(x_t) > \tau]$$

τ : validation-derived
low-FPR threshold
(reused unchanged)



**Optional temporal persistence
(three-of-five rule)**

Alert if ≥ 3 of the last 5 frames have $a_t = 1$.

Optional
post-processing
(validation selected)

No alert

Alert



**Forward to unchanged
victim VLM**



**External response policy,
not evaluated**